

# Outline

- RNNs and vanishing/exploding gradients
- Solutions



## RNNs: Advantages

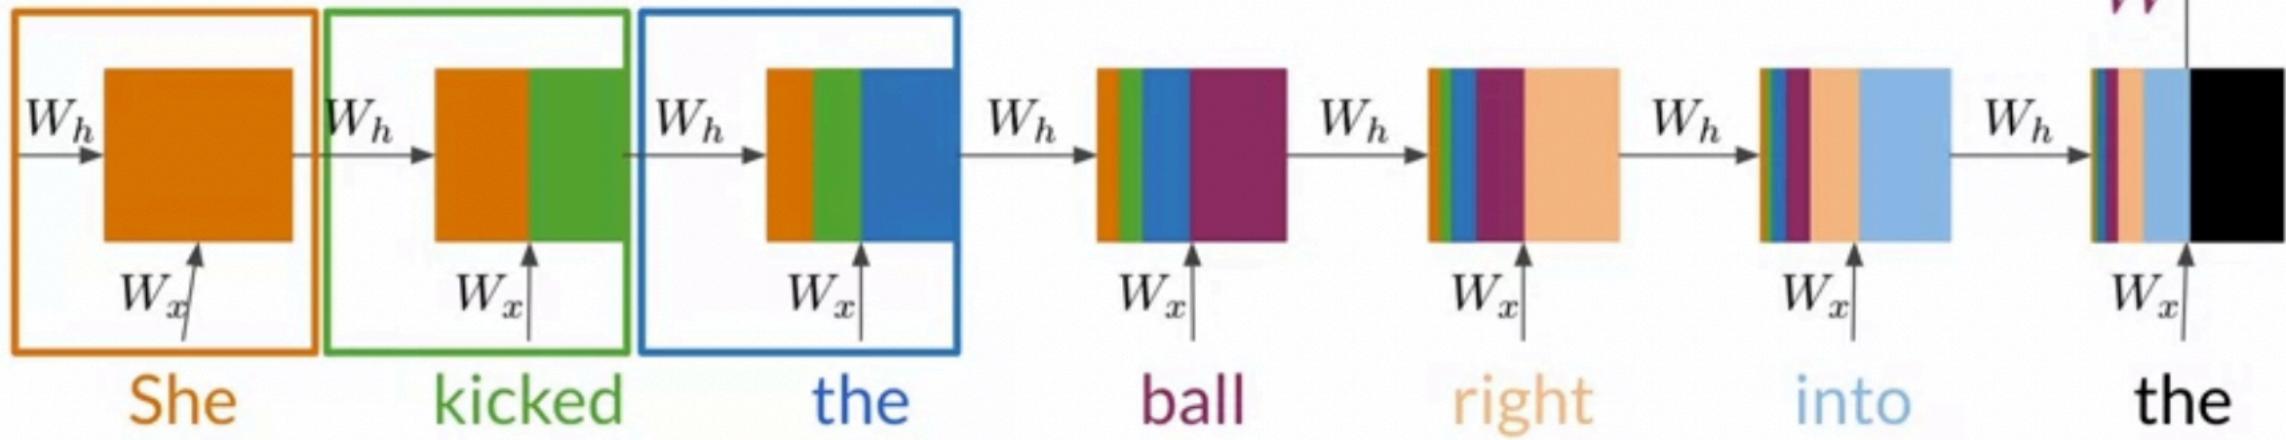
- + Captures dependencies within a short range
- + Takes up less RAM than other n-gram models

# RNNs: Disadvantages

- Struggles with longer sequences
- Prone to vanishing or exploding gradients

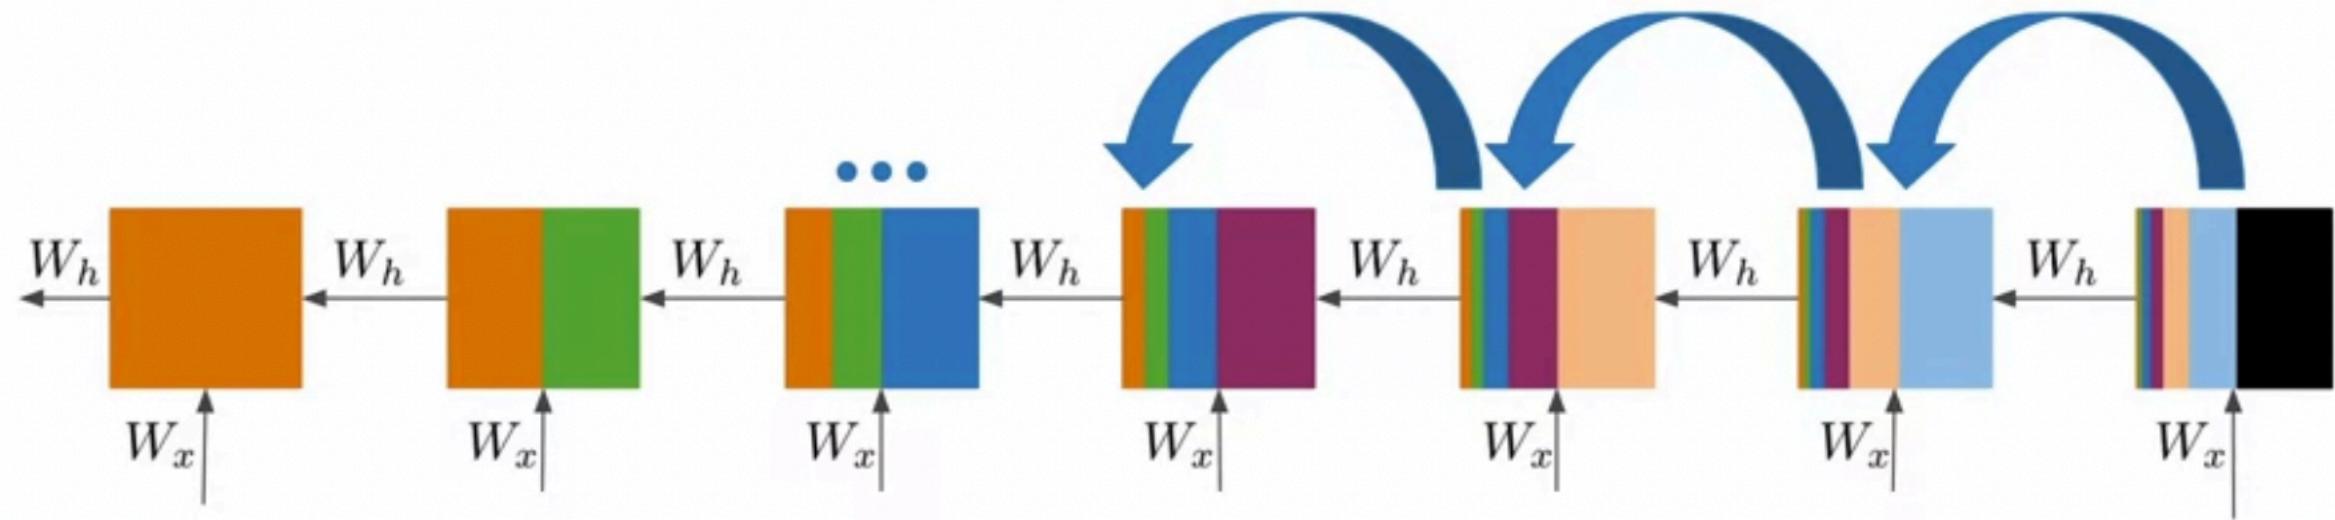
# RNN Basic Structure

She kicked the ball right into the \_\_\_\_\_



Learnable parameters

# Backpropagation through time



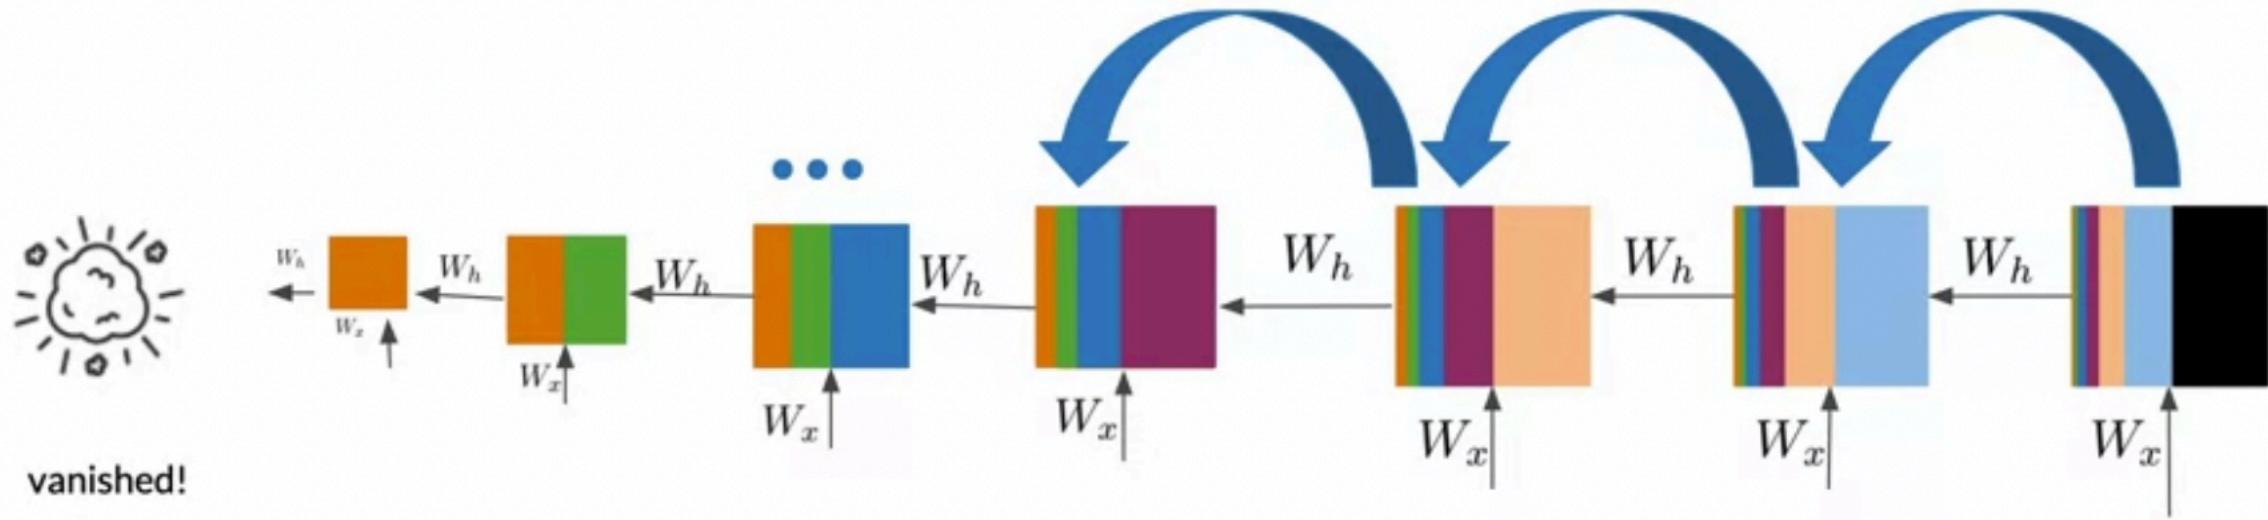
$$\int \sigma$$

between 0 and 1

$$\int \tanh$$

between -1 and 1

# The vanishing gradient problem



# Solving for vanishing or exploding gradients

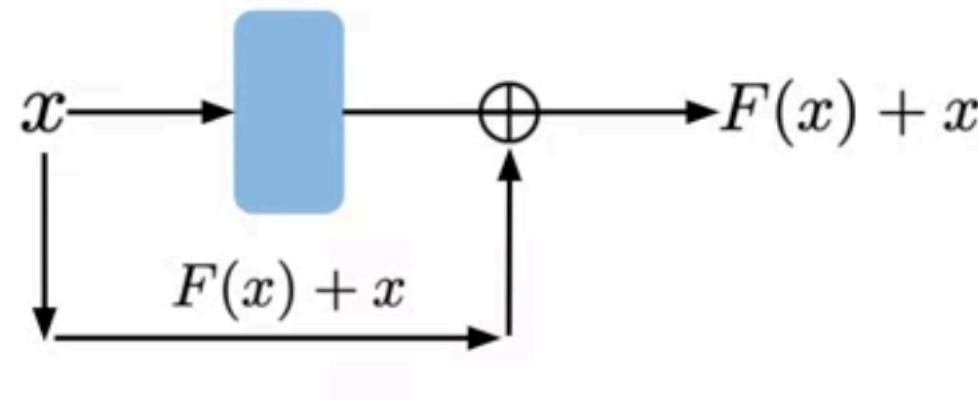
- Identity RNN with ReLU activation

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \quad -1 \longrightarrow 0$$

- Gradient clipping

$$32 \longrightarrow 25$$

- Skip connections



# Outline

- Meet the Long short-term memory unit!
- LSTM architecture
- Applications



# LSTMs: a memorable solution

- Learns when to remember and when to forget
- Basic anatomy:
  - A cell state
  - A hidden state with three gates
  - Loops back again at the end of each time step
- Gates allow gradients to flow unchanged

# LSTMs: Based on previous understanding

Cell state = before conversation

Forget gate = beginning of conversation

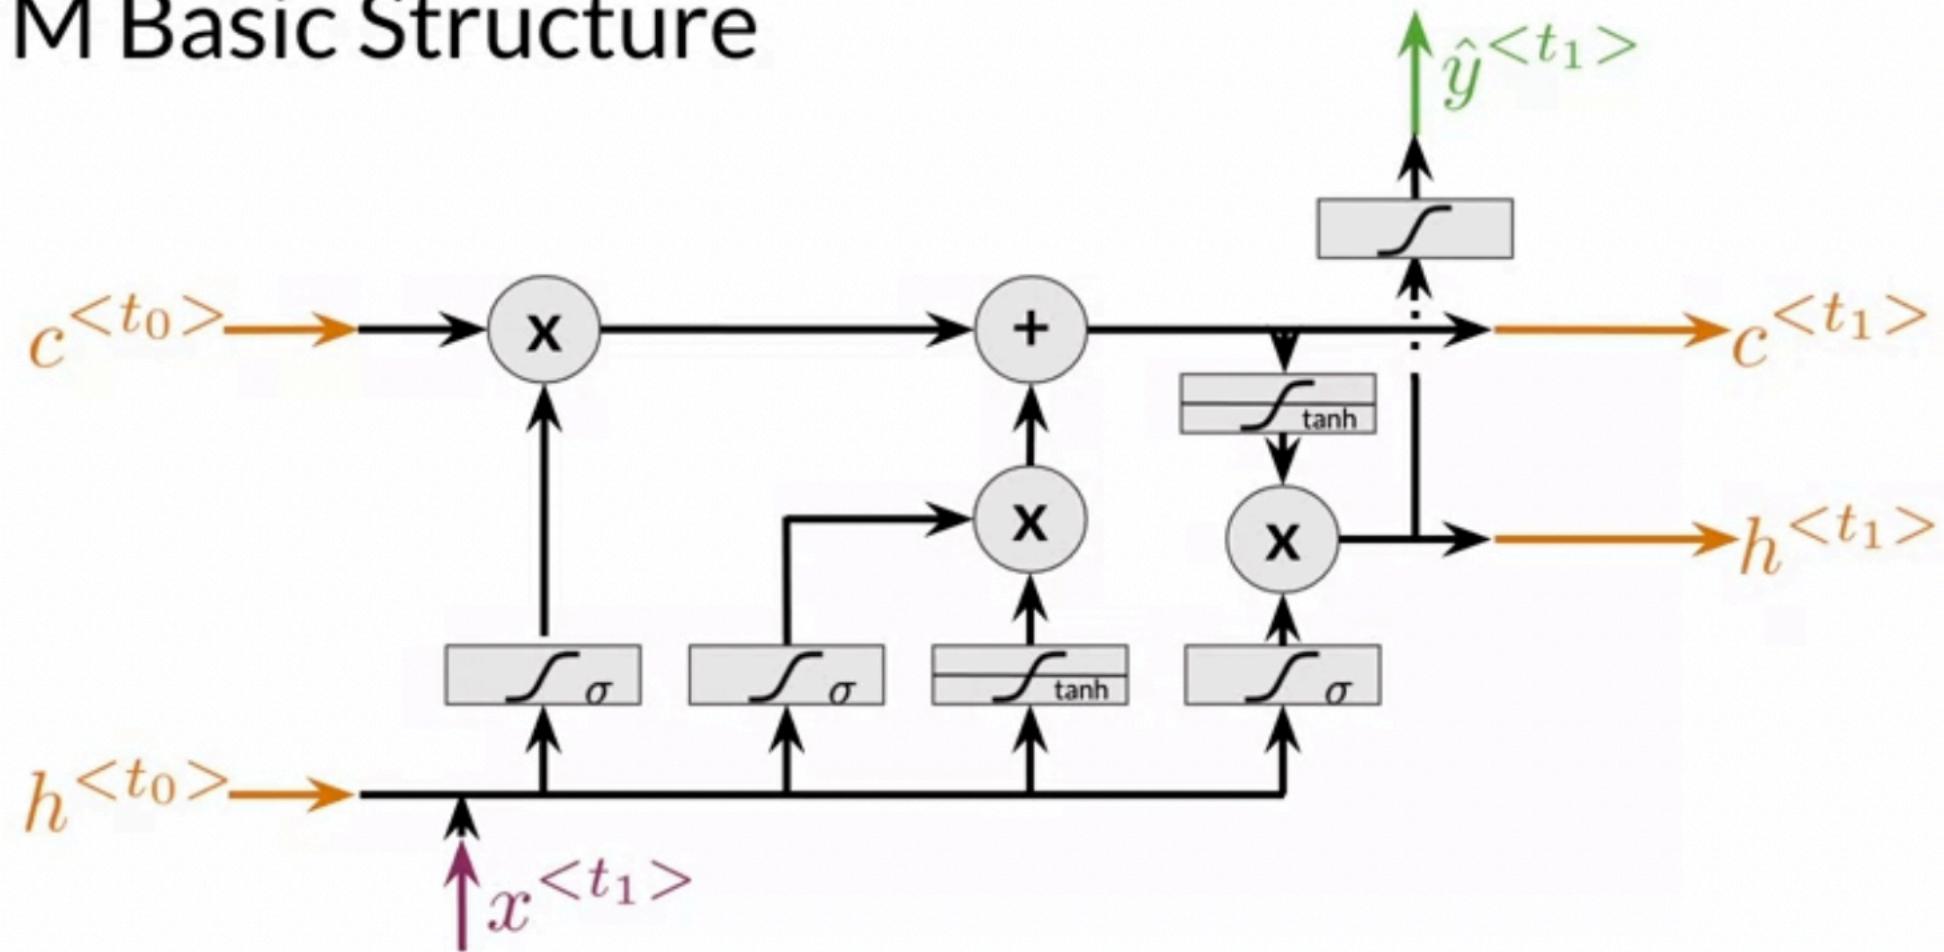
Input gate = thinking of a response

Output gate = responding

Updated cell state = after conversation



# LSTM Basic Structure



# Applications of LSTMs

Next-character  
prediction



Chatbots



Music  
composition

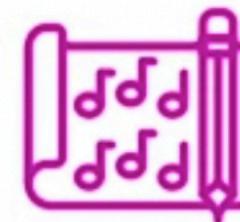


Image  
captioning

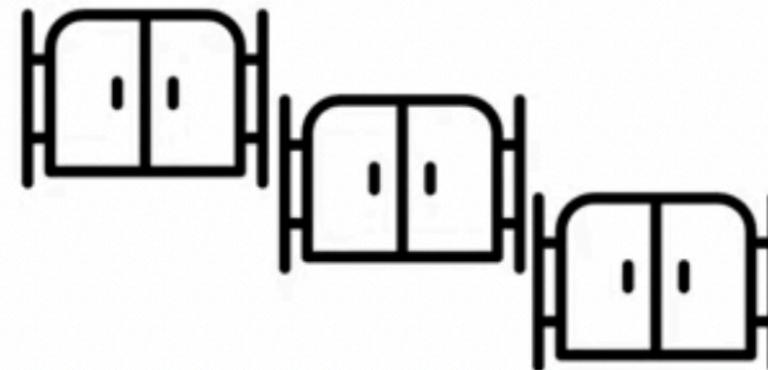


Speech  
recognition

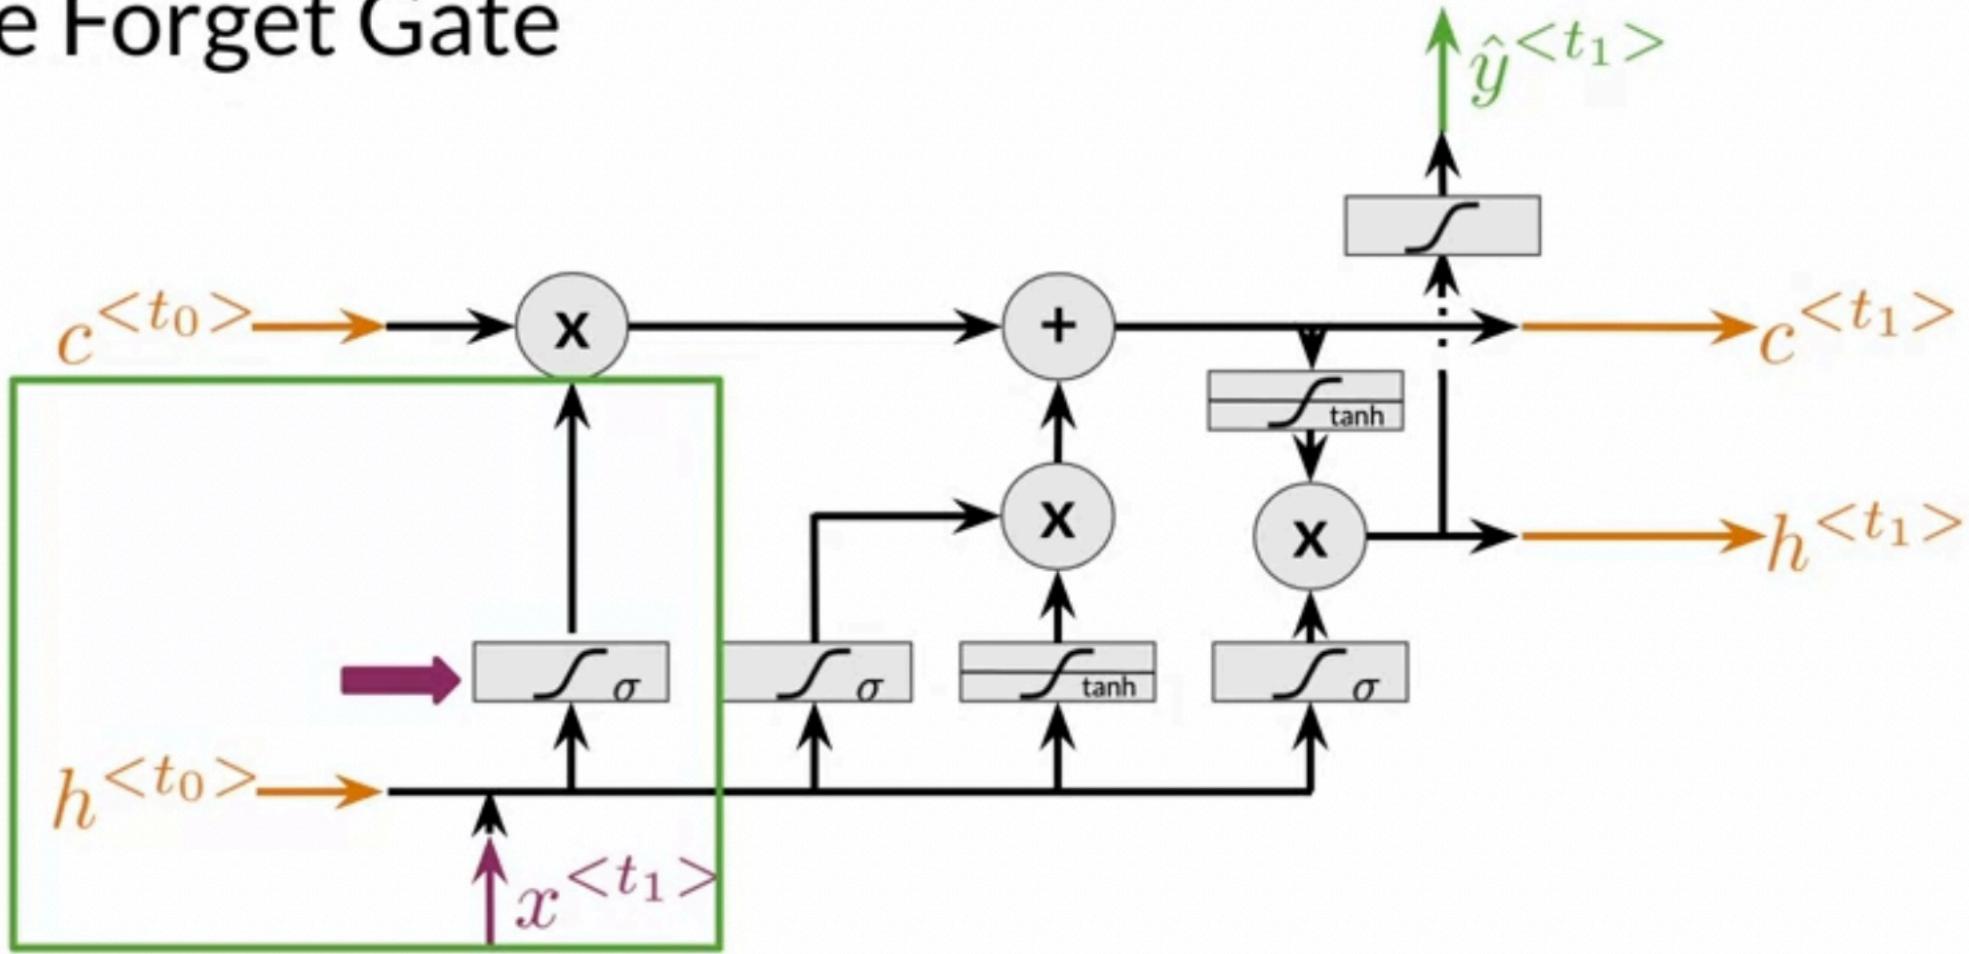


# Summary

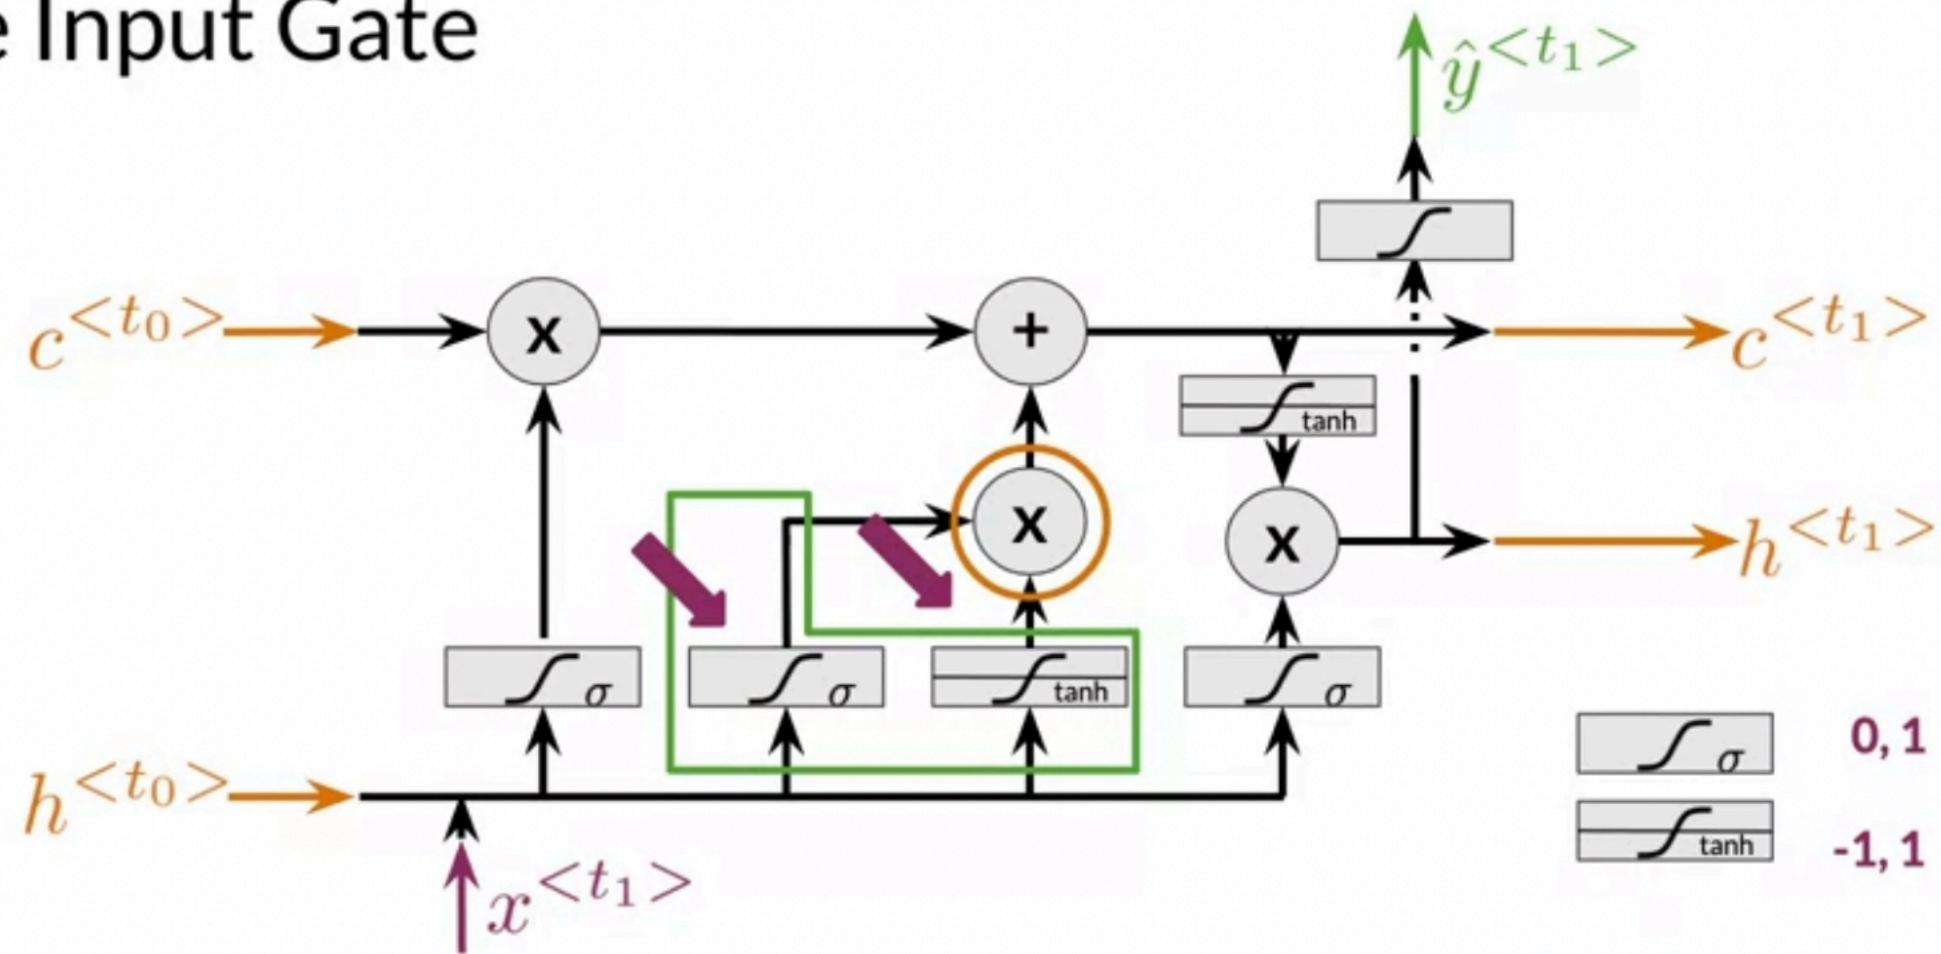
- LSTMs offer a solution to vanishing gradients
- Typical LSTMs have a cell and three gates:
  - Forget gate
  - Input gate
  - Output gate



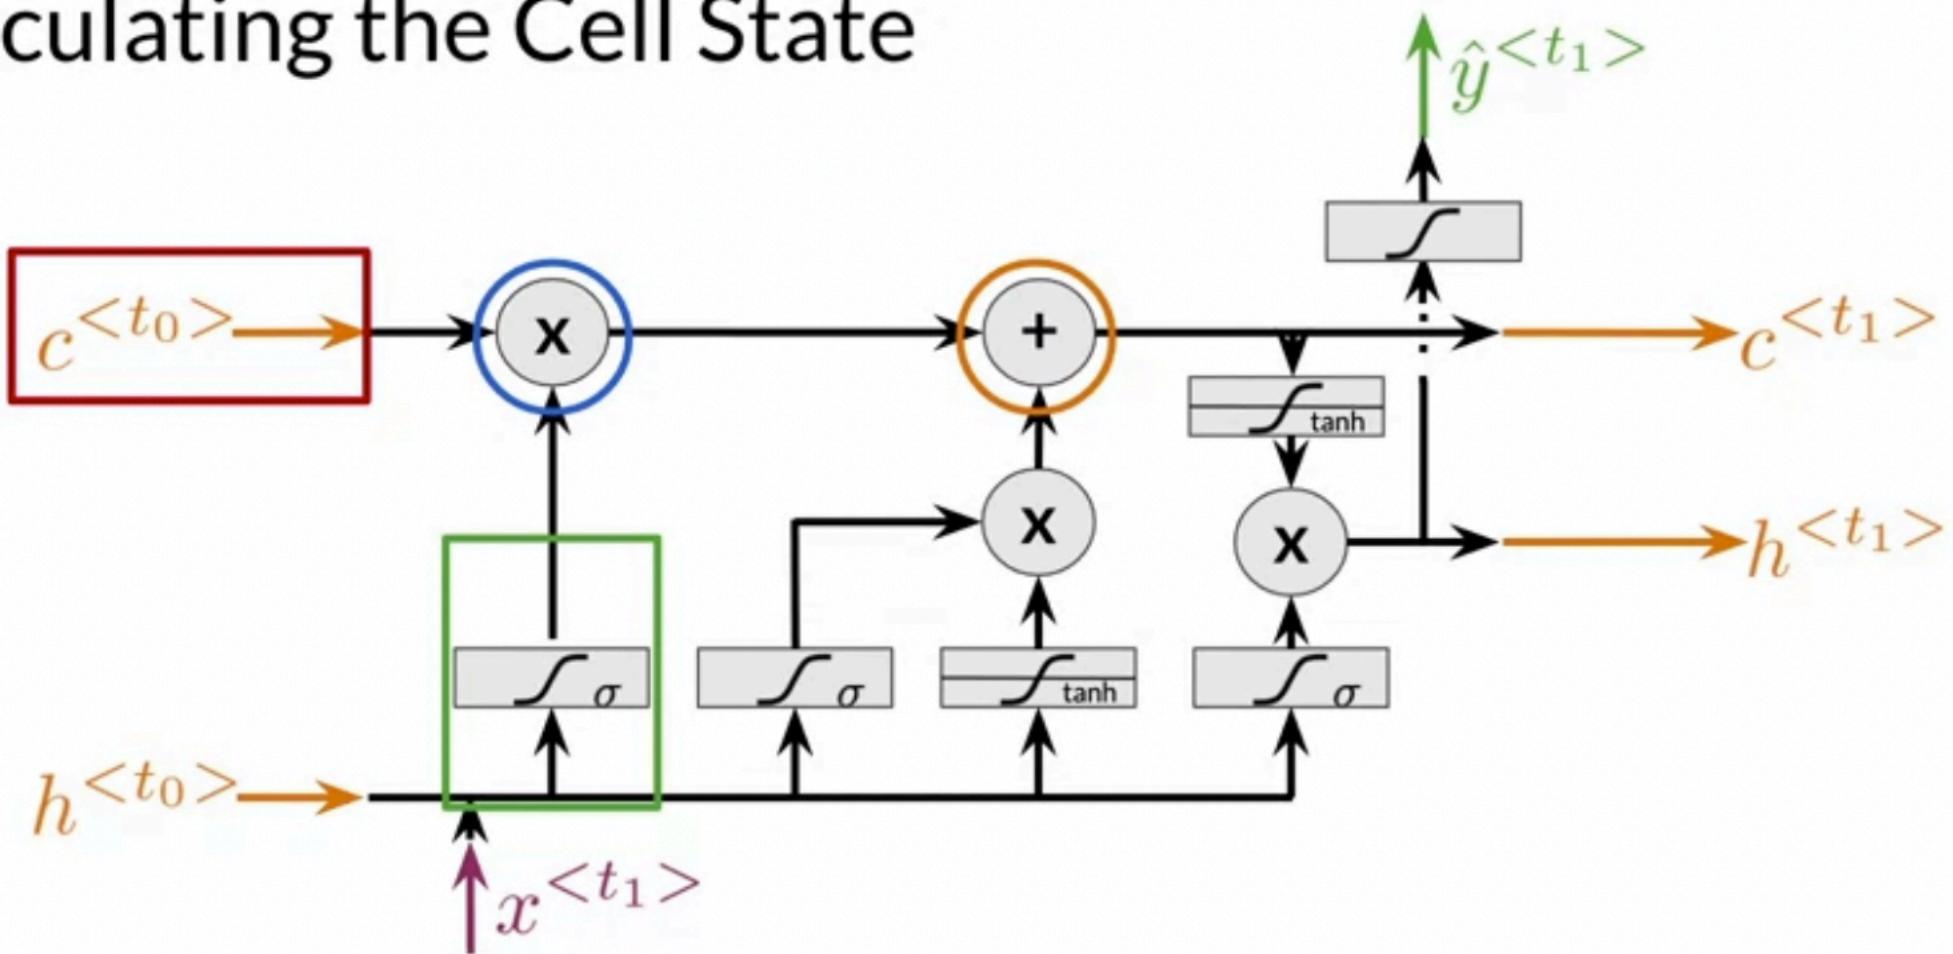
# The Forget Gate



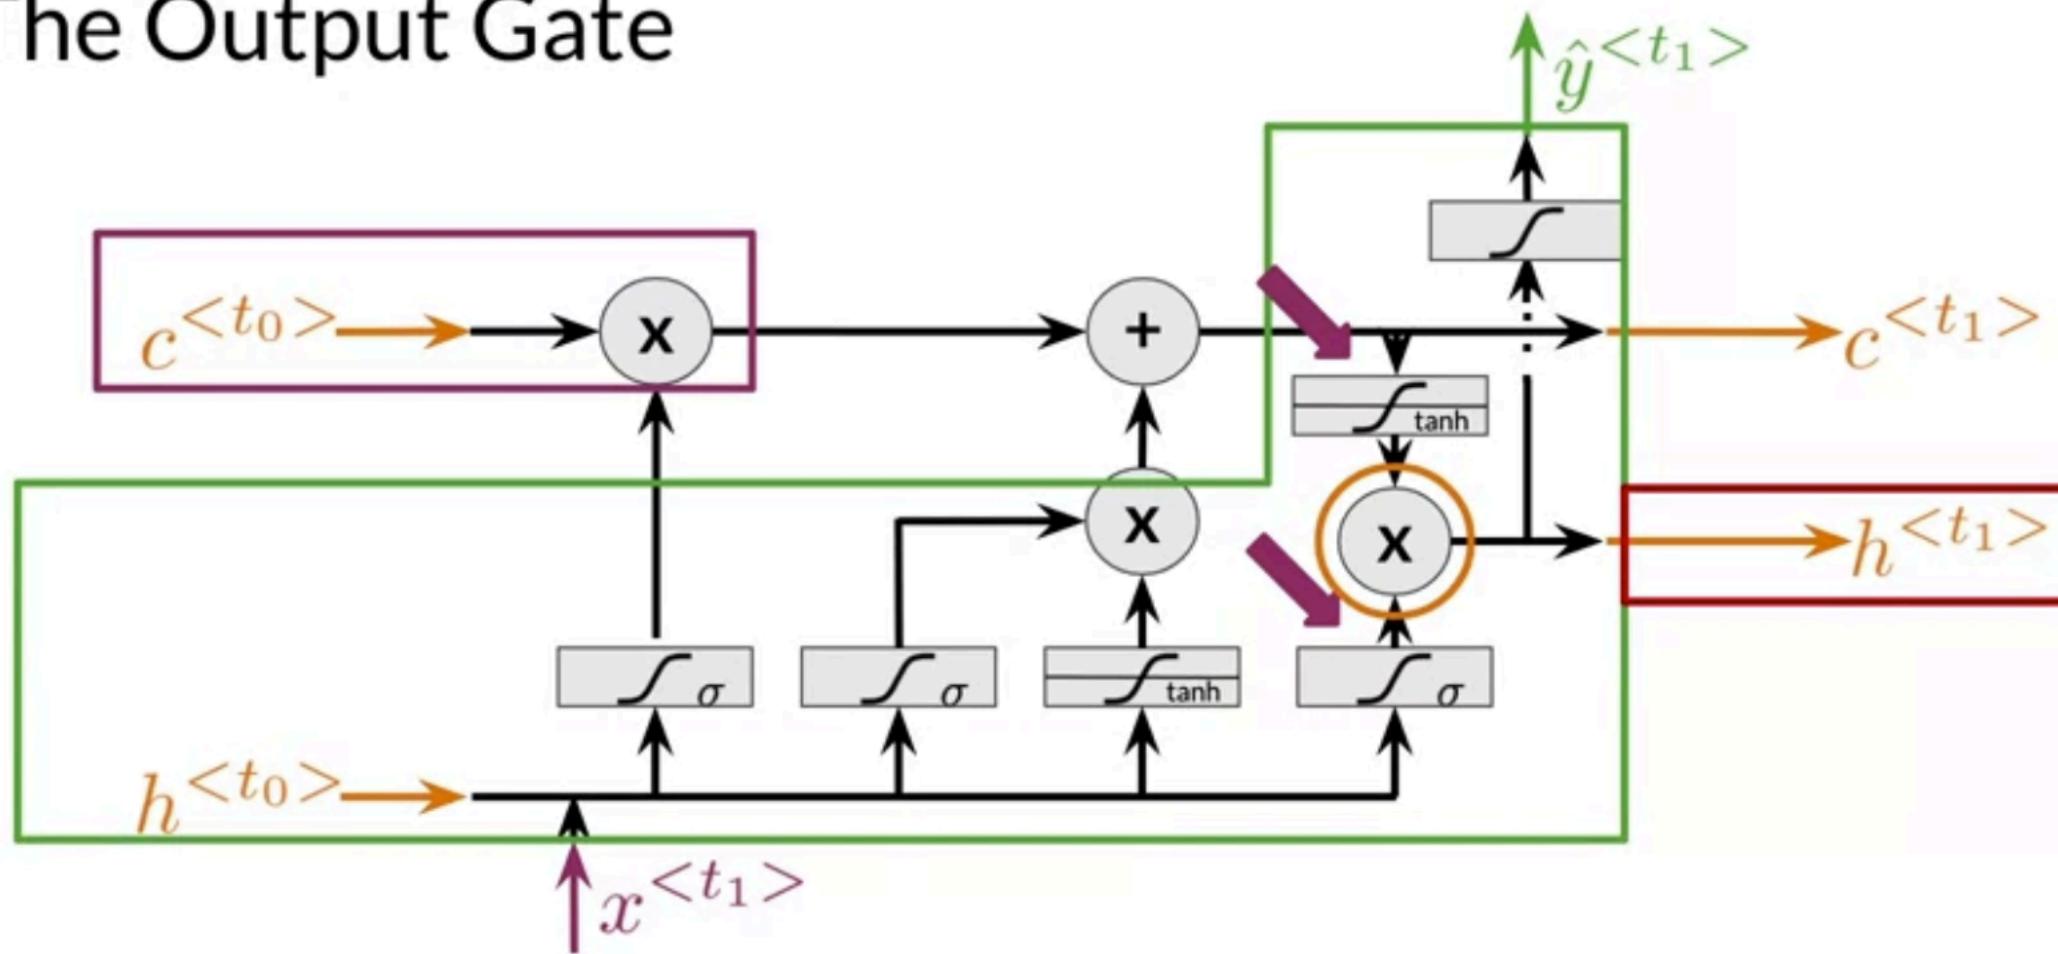
# The Input Gate



# Calculating the Cell State



# The Output Gate



# Summary

- LSTMs use a series of gates to decide which information to keep:
  - Forget gate decides what to keep
  - Input gate decides what to add
  - Output gate decides what the next hidden state will be
- One time step is completed after updating the states

# What is Named Entity Recognition?

- Locates and extracts predefined entities from text
- Places, organizations, names, time and dates



# Types of Entities



Thailand:  
Geographical

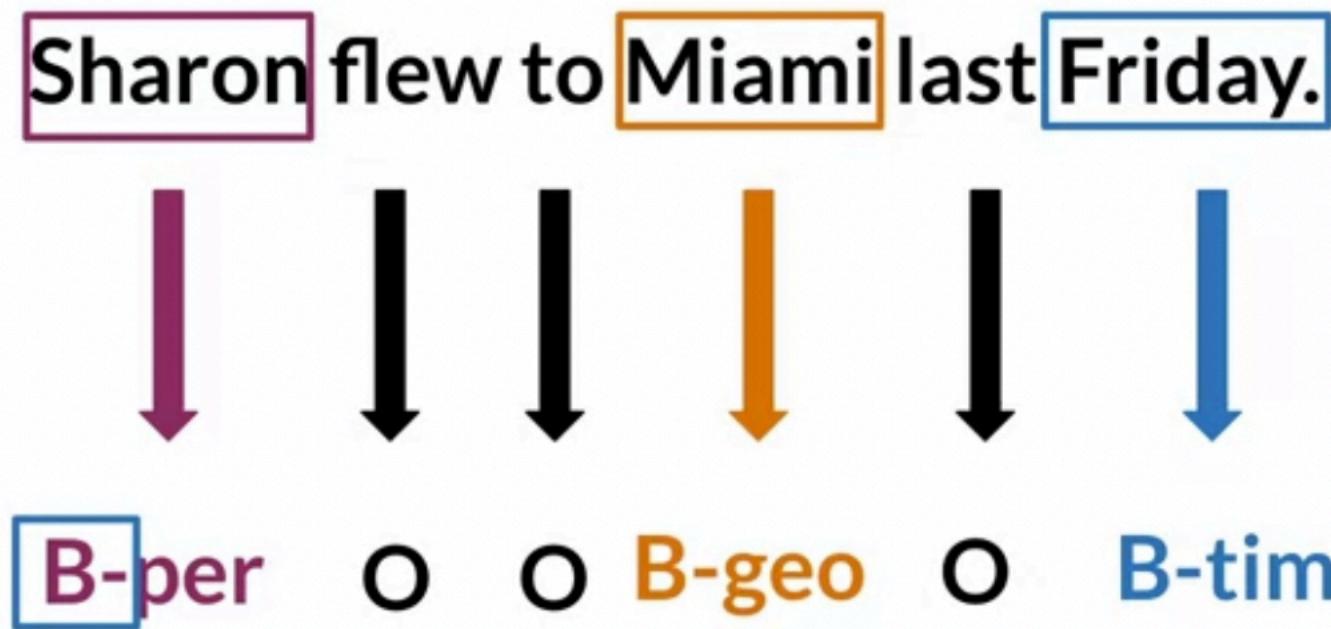


Google:  
Organization



Indian:  
Geopolitical

## Example of a labeled sentence



# Applications of NER systems

- Search engine efficiency
- Recommendation engines
- Customer service
- Automatic trading



# Outline

- Convert words and entity classes into arrays
- Token padding
- Create a data generator



# Processing data for NERs

- Assign each class a number

1

B-per

2

B-geo

3

B-tim

# Processing data for NERs

- Assign each class a number
- Assign each word a number

Sharon flew to Miami last Friday.

[ 4282, 853, 187, 5388, 2894, 7 ]

B-per O O B-geo O B-tim

# Token padding

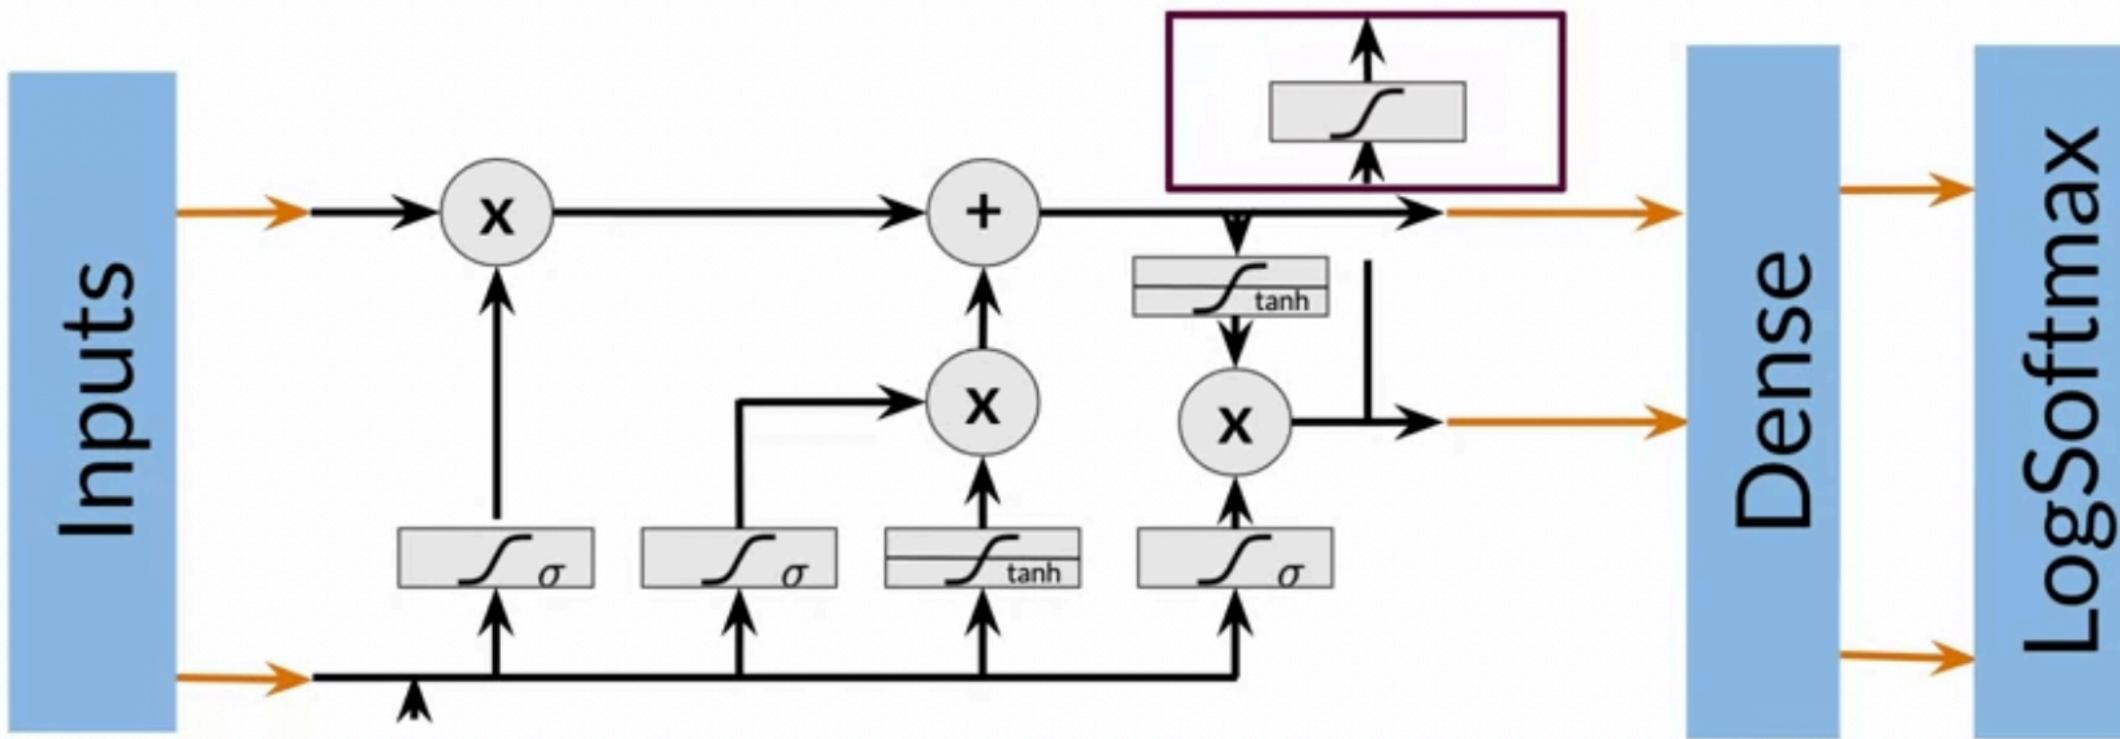
For LSTMs, all sequences need to be the same size.

- Set sequence length to a certain number
- Use the <PAD> token to fill empty spaces

# Training the NER

1. Create a tensor for each input and its corresponding number
2. Put them in a batch  64, 128, 256, 512 ...
3. Feed it into an LSTM unit
4. Run the output through a dense layer
5. Predict using a log softmax over K classes

# Training the NER



# Layers in Trax

```
model = tl.Serial(  
    tl.Embedding(),  
    tl.LSTM(),  
    tl.Dense()  
    tl.LogSoftmax()  
)
```

# Summary

- Convert words and entities into same-length numerical arrays
- Train in batches for faster processing
- Run the output through a final layer and activation



# Evaluating the model

1. Pass test set through the model
2. Get arg max across the prediction array
3. Mask padded tokens
4. Compare outputs against test labels

# Evaluating the model in Python

```
def evaluate_model(test_sentences, test_labels, model):
    pred = model(test_sentences)
    outputs = np.argmax(pred, axis=2)
    mask = ...
    accuracy = np.sum(outputs==test_labels)/float(np.sum(mask))

    return accuracy
```

# Summary

- If padding tokens, remember to mask them when computing accuracy
- Coding assignment!